Introduction to Big Data
Graded Assignment

Name - Avijeet Palit
Roll - 21f1005675
Date - 22/11/2024

Objective

Convert the Spark MLib code from the Databricks decision trees example to use the CrossValidator autotuner. Analyze and report the best-performing model parameters.

Detailed Steps

Step 1: Create a GCS Bucket

- Navigate to the Cloud Storage Section:
- Log in to your Google Cloud Platform (GCP) Console.
- Click on Cloud Storage from the navigation menu.
- Create a New Bucket:

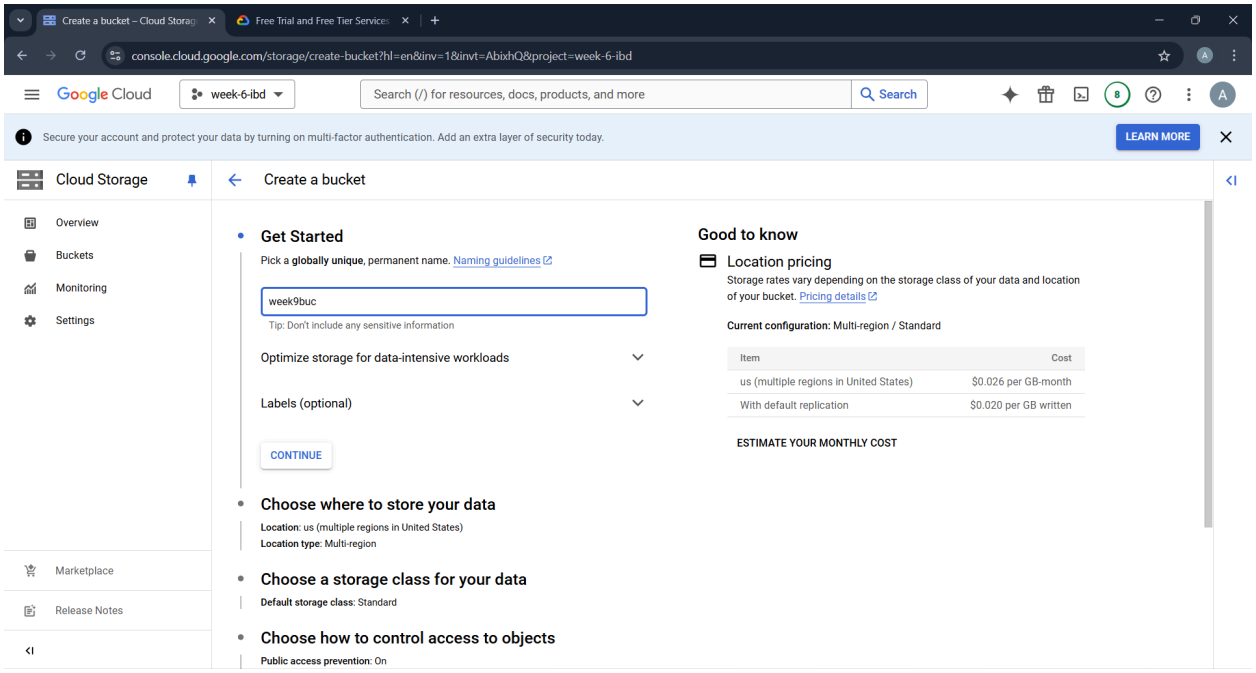

- Provide a name for the bucket (e.g., "week9buc").
- Set Location and Permissions:
- Select the desired bucket location (e.g., Regional/Multiregional).

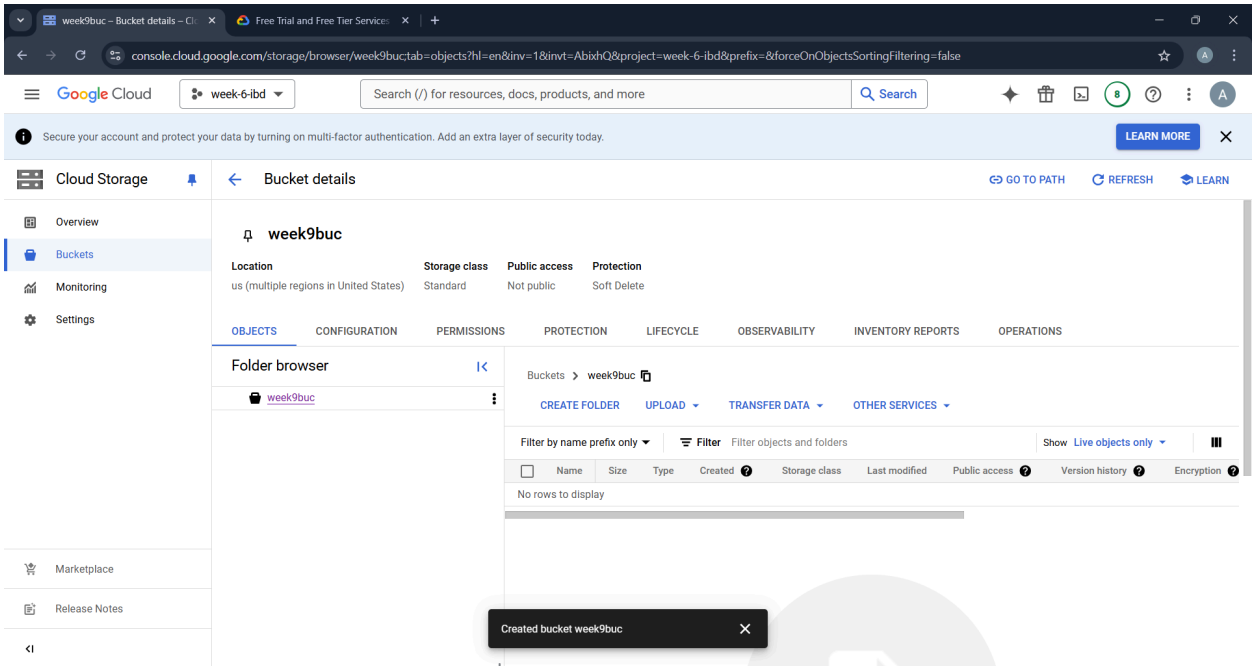

Step 2: Generate or Use MNIST Data

- Obtain MNIST Dataset:
- Either generate the MNIST Train and Test datasets or download them from a reliable source.

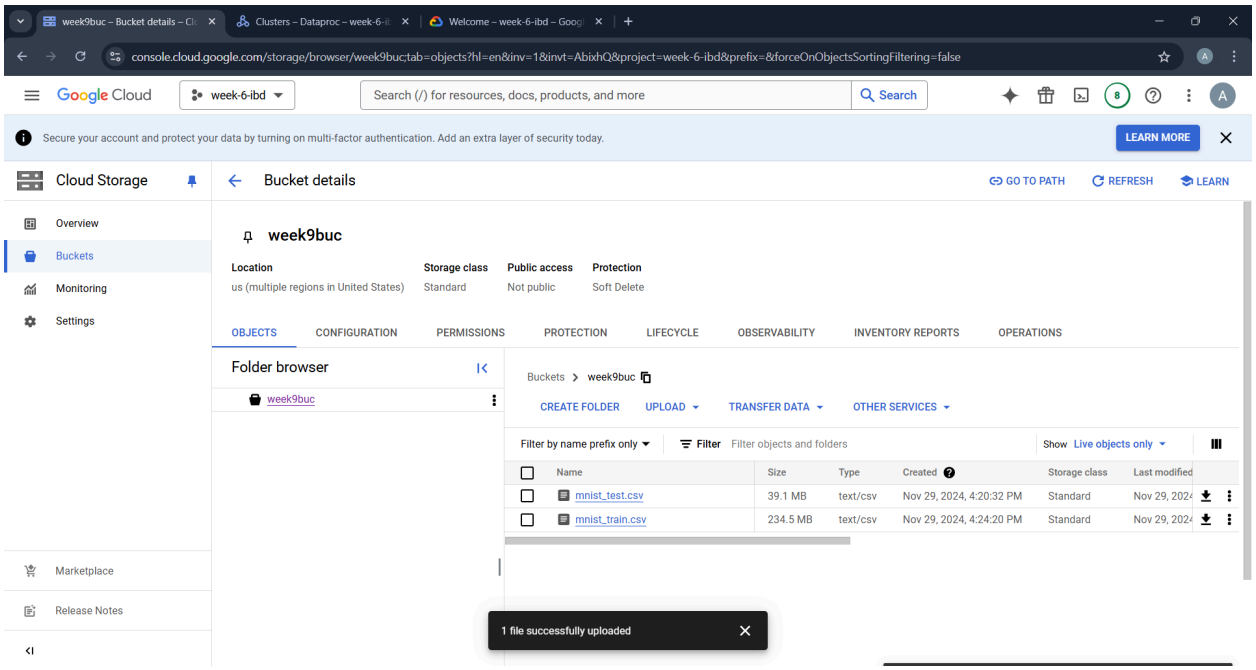

```python
import numpy as np
import pandas as pd
from tensorflow.keras.datasets import mnist

# Step 1: Load the MNIST dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()

# Step 2: Flatten the 28x28 images into a single 784-length vector
X_train_flat = X_train.reshape(X_train.shape[0], -1)
X_test_flat = X_test.reshape(X_test.shape[0], -1)

# Step 3: Normalize the data to [0, 1]
X_train_flat = X_train_flat.astype('float32') / 255.0
X_test_flat = X_test_flat.astype('float32') / 255.0

# Step 4: Create Pandas DataFrames
# Combine the features and labels for both train and test datasets
train_data = pd.DataFrame(X_train_flat)
train_data['label'] = y_train

test_data = pd.DataFrame(X_test_flat)
test_data['label'] = y_test

# Step 5: Save the datasets to CSV files
train_data.to_csv('mnist_train.csv', index=False)
test_data.to_csv('mnist_test.csv', index=False)

print("MNIST Train and Test datasets saved as 'mnist_train.csv' and 'mnist_test.csv'")
```

- Upload Data to the GCS Bucket:

Step 3: Create a Dataproc Cluster

- Go to the Dataproc Section:
- In the GCP Console, navigate to the Dataproc page.
- Create a Cluster:
- Click on Create Cluster.
- Configure Cluster Settings:
    - Set Cluster Type to "Standard".
    - Choose appropriate machine types for the master and worker nodes based on your data size and computational needs.
    - Enable the Component Gateway for easier job monitoring and debugging.
    - Ensure Spark Compatibility

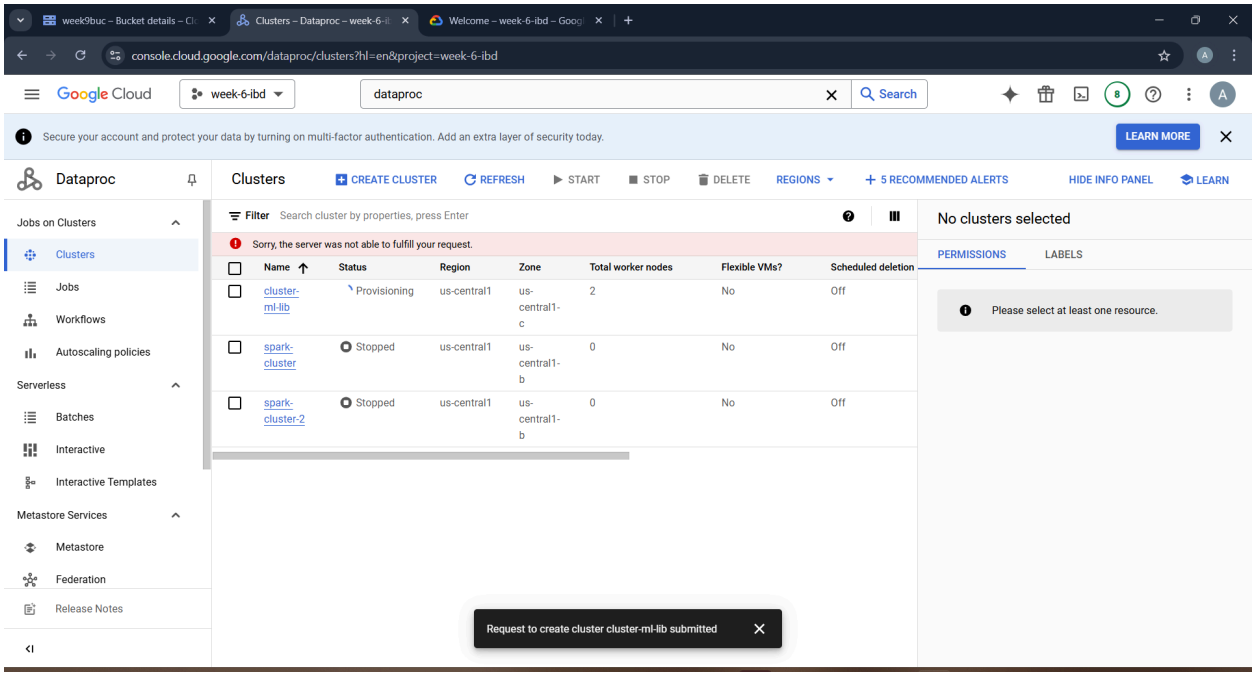

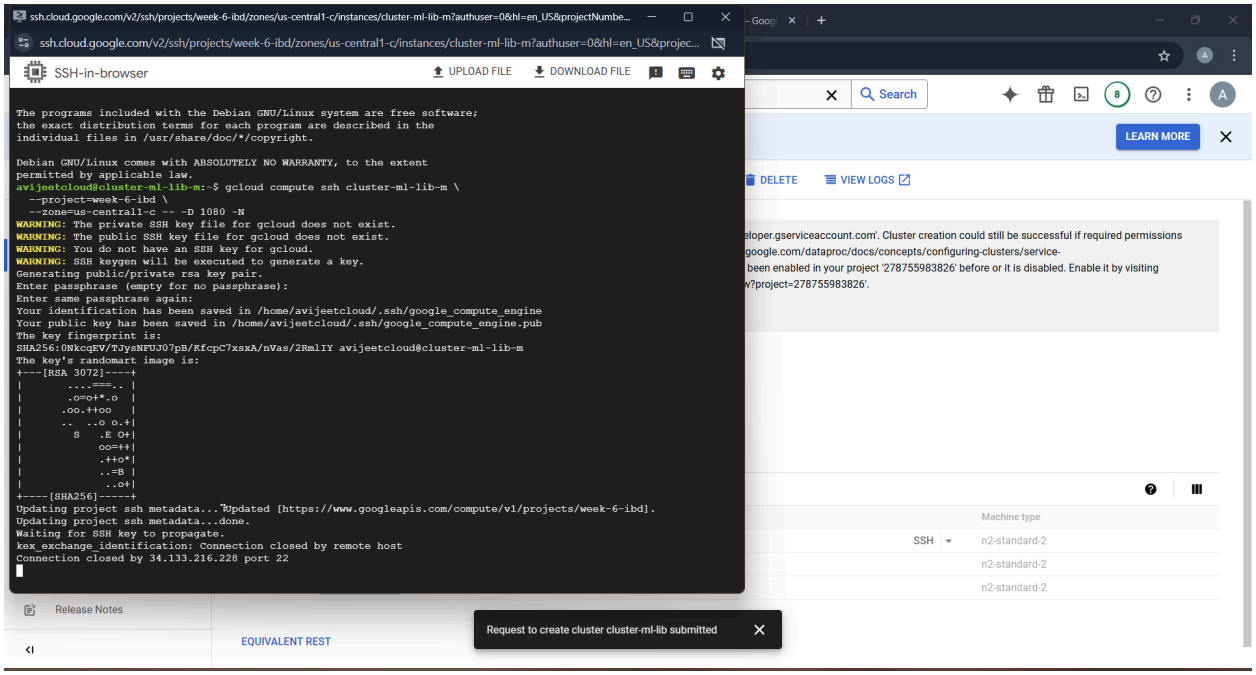

Step 4: Write and Modify the Spark Code

- Fetch Base Code:
  - Download the sample decision trees code from the Databricks example provided in the assignment.

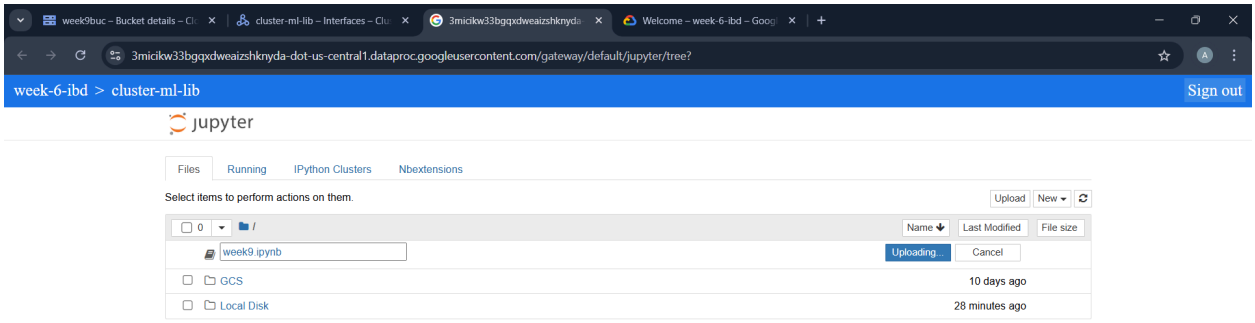

- Modify the Code to Use CrossValidator:
  - Use CrossValidator to tune the model parameters (e.g., maxDepth, maxBins).
  - Update Model Evaluation Metrics:
  - Adjust the code to output performance metrics for each model.

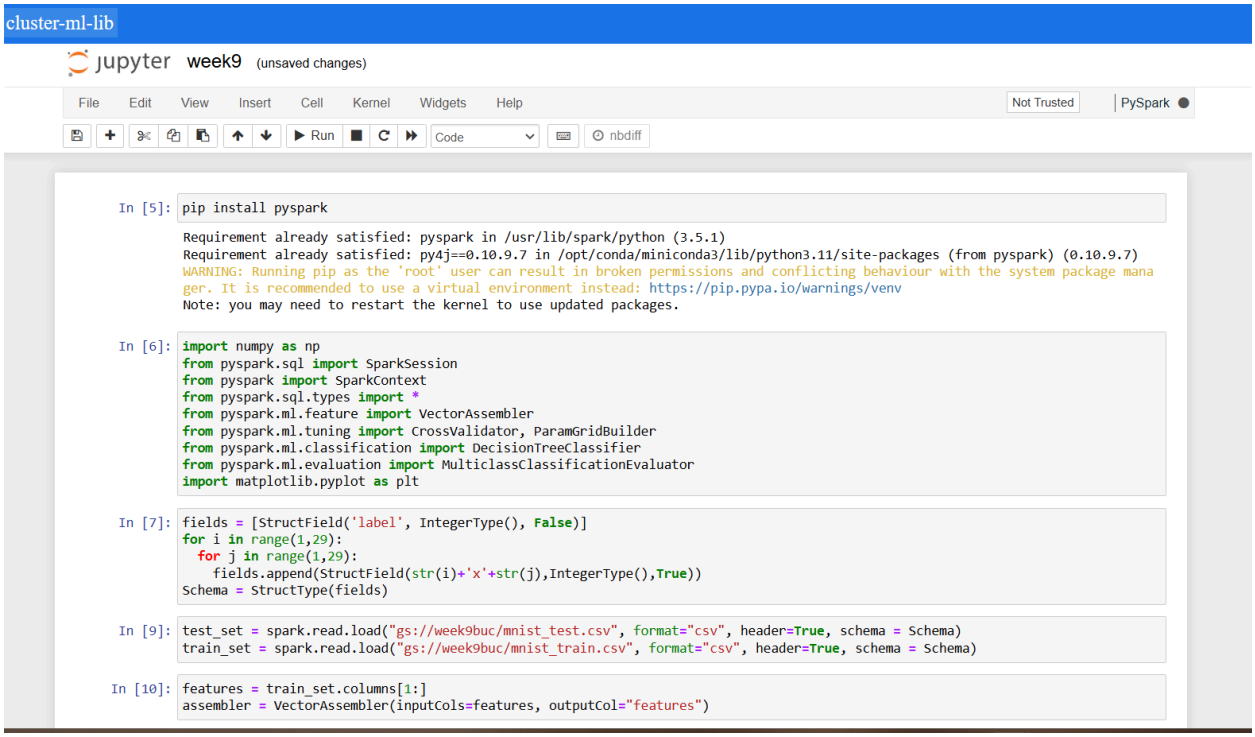

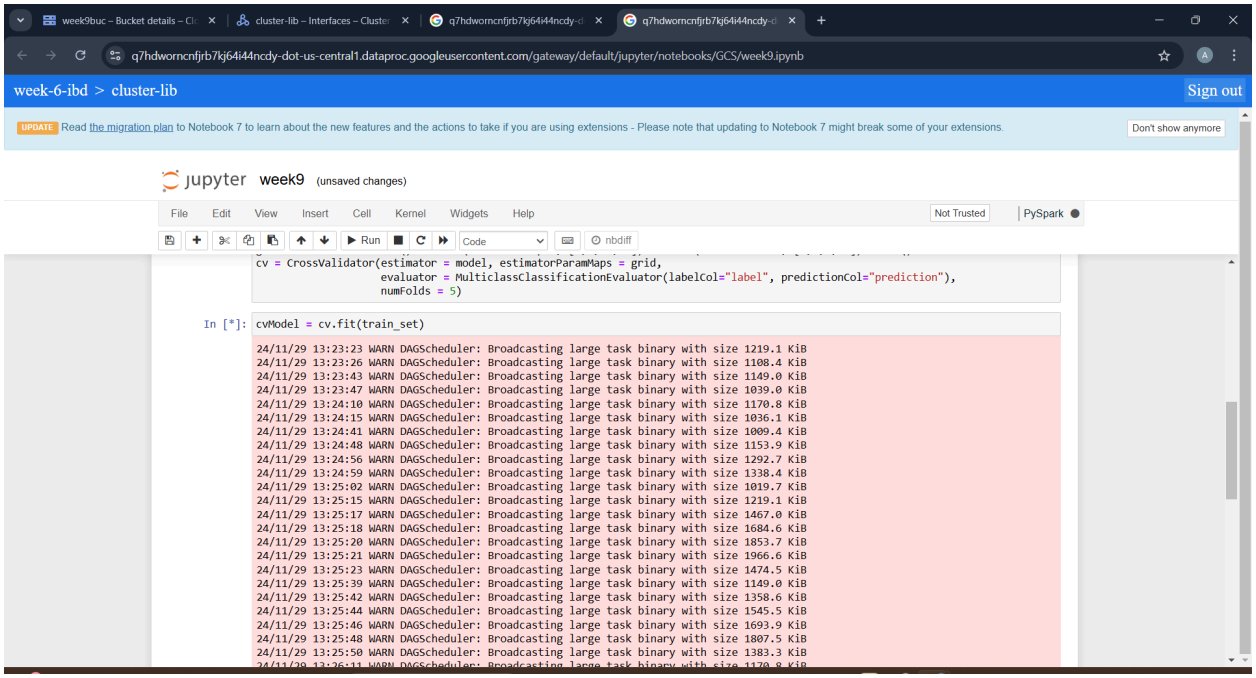

Results

- Best Parameters after CrossValidator Autotuning:
- List the parameters that yielded the best performance (e.g., maxDepth = 10, maxBins = 32).

In [9]:
```
est = cvModel.bestModel
print(f'Best tree depth {est.getMaxDepth()}, Best bin size {est.getMaxBins()}')
```

```
Best tree depth 16, Best bin size 2
```

In [10]:
```
testresult = est.transform(test_set)
```

In [11]:
```
print(f'Total test points {testresult.count()}')
wrong = testresult.where(testresult.label!=testresult.prediction)
print(f'Wrongly classified count {wrong.count()}')
```

```
Total test points 9999
```

```
24/11/29 14:15:48 WARN DAGScheduler: Broadcasting large task binary with size 1606.9 KiB
[Stage 544:============================>                            (1 + 1) / 2]
```

```
Wrongly classified count 1124
```

- Misclassified Data Points:
- Analyze the misclassified data points to identify potential patterns or weaknesses in the model.

In [12]:
```
temp = wrong.limit(5).select("label","prediction","features").collect()

for r in temp:
    plt.figure()
    row = np.array(r.features).reshape((28,28))
    plt.imshow(row,cmap = 'gray')
    plt.title(f'True {r.label} Predicted {r.prediction}')
```